

Carnap's Robot Redux: LLMs, Intensional Semantics, and the Implementation Problem in Conceptual Engineering

Bradley P. Allen

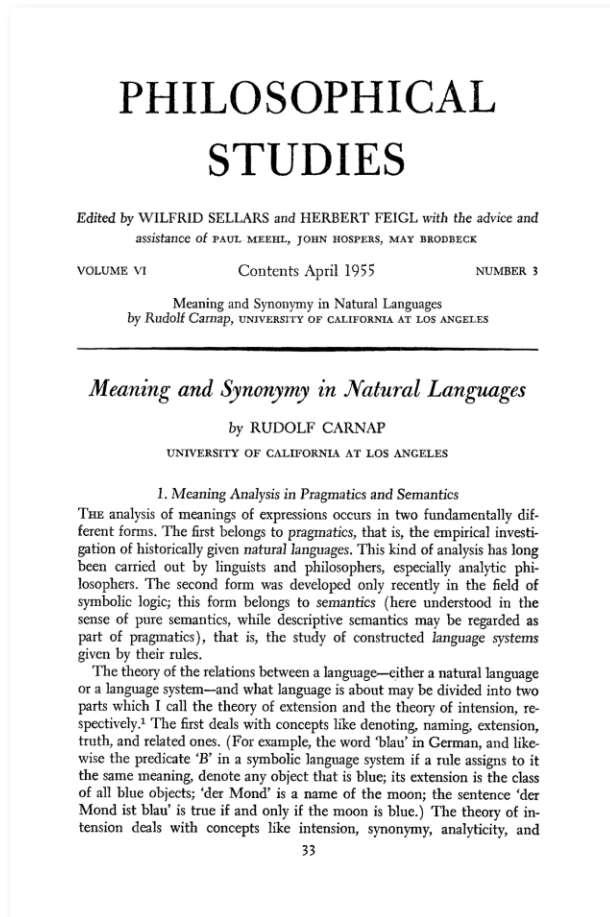
LLMs and Philosophy, Kanazawa

2024-09-27

Overview

- Carnap's Robot Redux
 - (Carnap 1955) through the lens of LLMs and mechanistic interpretability
 - Intensional semantics based on zero-shot chain-of-thought classifiers
- Experiments in conceptual engineering and knowledge engineering
 - Intentions-as-classifiers as targets of conceptual engineering
 - Intentions-as-classifiers for the evaluation of knowledge graphs
- LLMs as cognitive tools for philosophy
 - For experimental conceptual engineering
 - For ameliorative analysis of knowledge bases as a solution for the implementation problem in conceptual engineering
- Future work on LLMs and meaning
 - LLMs, metalinguistic disagreement, hyperintensionality, and two-component semantics

(Carnap 1955)



- "Guided by the view that “the assignment of an intension [i.e. meaning] is an empirical hypothesis which...can be tested by observations of language behavior” (Carnap 1955, 37), Carnap described a method for uncovering intensions that involves presenting language users with a range of logically possible scenarios and asking them to make judgments regarding the concept in question." (Shepherd and Justus 2015)

Carnap's Robot



"the investigator of X's language L"

"a general description of a kind of object"

"a predicate"

"affirmation, denial, or abstention"



"robot X"

Could an LLM play the role of Carnap's Robot?

- In recent years, large language models (LLMs) have emerged as a technology that promises to be of “substantial value in the scientific study of language learning and processing” (Mahowald et al 2023)
- Given this, we ask the question: might LLMs be useful in the conduct of investigations of the sort Carnap proposed in (Carnap 1955)?
- We argue that that is the case
- We start by showing how an LLM can be used to implement intensions

Intensions as classifiers

- (Muskens 2005): intensions as logic programs
- (Larsson 2015): intensions as perceptron-based classifiers
 - "The crucial step in making use of classifiers in formal semantics is to regard them as (parts of) representations of intensions of linguistic expressions."
- Our work: intensions as zero-shot chain-of-thought classifiers



Token sequences, LLMs, and prompts

- Let Σ be a countable set of tokens and Σ^* be the set of finite sequences of tokens $\lceil t_0 \dots t_k \rceil$, where $t_i \in V$, $0 \leq i \leq k$, $k \in \mathbb{N}$.
- For $P, x_0, \dots, x_n, y_0, \dots, y_n \in \Sigma^*$, let $P_{y_0, \dots, y_n}^{x_0, \dots, x_n}$ be the sequence of tokens obtained from P by replacing x_i , wherever it occurs in P , by y_i , where $0 \leq i \leq n$, $n \in \mathbb{N}$.
- Given a corpus $C \in \Sigma^*$, let L_C be a an LLM trained on a sample of C .
- Let $\Pi : L_C \times \Sigma^* \rightarrow \Sigma^*$ be a function such that $\Pi(L_C, P)$ is the output sequence generated by the language model given the input sequence, or prompt, P .

Predicative language

We define a formal language \mathcal{L} built from constants, variables, binary relation symbols, the Boolean connectives \neg , \wedge , and \vee , and round parentheses as auxiliary symbols.

- If R is a binary relation symbol and a and b are terms, then aRb is an atomic formula.
- Let $\mathcal{L}_{AT} \subset \Sigma^*$ be the set of atomic formulas.
- The set of terms of \mathcal{L} are the constants and variables in \mathcal{L} .
- The set of formulas of \mathcal{L} are the items in \mathcal{L}_{AT} , and if ϕ and ψ are formulas and x is a variable, so are the following:

$$\neg \phi \mid (\phi \wedge \psi) \mid (\phi \vee \psi) \mid (\phi \rightarrow \psi) \mid (\phi \leftrightarrow \psi) \mid \exists x \phi \mid \forall x \phi$$

Models

- A model $\mathcal{M} = (D, S, L_C, P)$ for \mathcal{L} is a tuple where:
 - D is a countable set of individuals
 - $S \subset \mathcal{P}(\mathcal{L}_{AT})$ is a countable set of states
 - L_C is a large language model trained on a corpus C
 - P be a prompt that instructs L_C to take an atomic formula ϕ and a state σ and generate the token sequence $\ulcorner 1 \urcorner$ if ϕ is true given the state σ
- An assignment g is a function that assigns to each variable an element of D .

Intensions

- The intension of a term t with respect to a model \mathcal{M} and an assignment g , denoted by $\llbracket t \rrbracket^{\mathcal{M},g}$, is defined as follows:
 - If t is a constant, then $\llbracket t \rrbracket^{\mathcal{M},g}(s) = d \in D$ for all $s \in S$.
 - If t is a variable, then $\llbracket t \rrbracket^{\mathcal{M},g}(s) = g(t)$ for all $s \in S$.
- The intension of a binary relation symbol R with respect to a model \mathcal{M} and an assignment g , denoted by $\llbracket R \rrbracket^{\mathcal{M},g}$, is a function from S to $\mathcal{P}(D \times D)$, defined as follows:
 - For each $s \in S$ and $d_1, d_2 \in D$, $(d_1, d_2) \in \llbracket R \rrbracket^{\mathcal{M},g}(s)$ iff $\Pi(L_C, P_{d_1 R d_2, s}^{\phi, \sigma}) = \ulcorner 1 \urcorner$.

Extensions

- The extension of a formula ϕ with respect to a model \mathcal{M} , an assignment g , and a state s , denoted by $\llbracket \phi \rrbracket^{\mathcal{M},g,s}$, is a truth value, defined recursively as follows:
 - If ϕ is an atomic formula of the form aRb , then $\llbracket \phi \rrbracket^{\mathcal{M},g,s} = 1$ if and only if $(\llbracket a \rrbracket^{\mathcal{M},g}(s), \llbracket b \rrbracket^{\mathcal{M},g}(s)) \in \llbracket R \rrbracket^{\mathcal{M},g}(s)$, otherwise 0.
 - The extension of complex formulas is defined as usual in classical logic, using the truth values of their subformulas.

Knowledge graphs as objects of intensional investigations

- Now that we have a hammer for investigating intensions, we need some nails
 - A *knowledge graph* represents knowledge using nodes for entities and edges for relations (Hogan et al. 2021)
 - Knowledge graphs are key information infrastructure for many Web applications (Heist et al. 2020)
 - Following (Angles et al. 2020), we use the RDF data model to describe knowledge graphs.

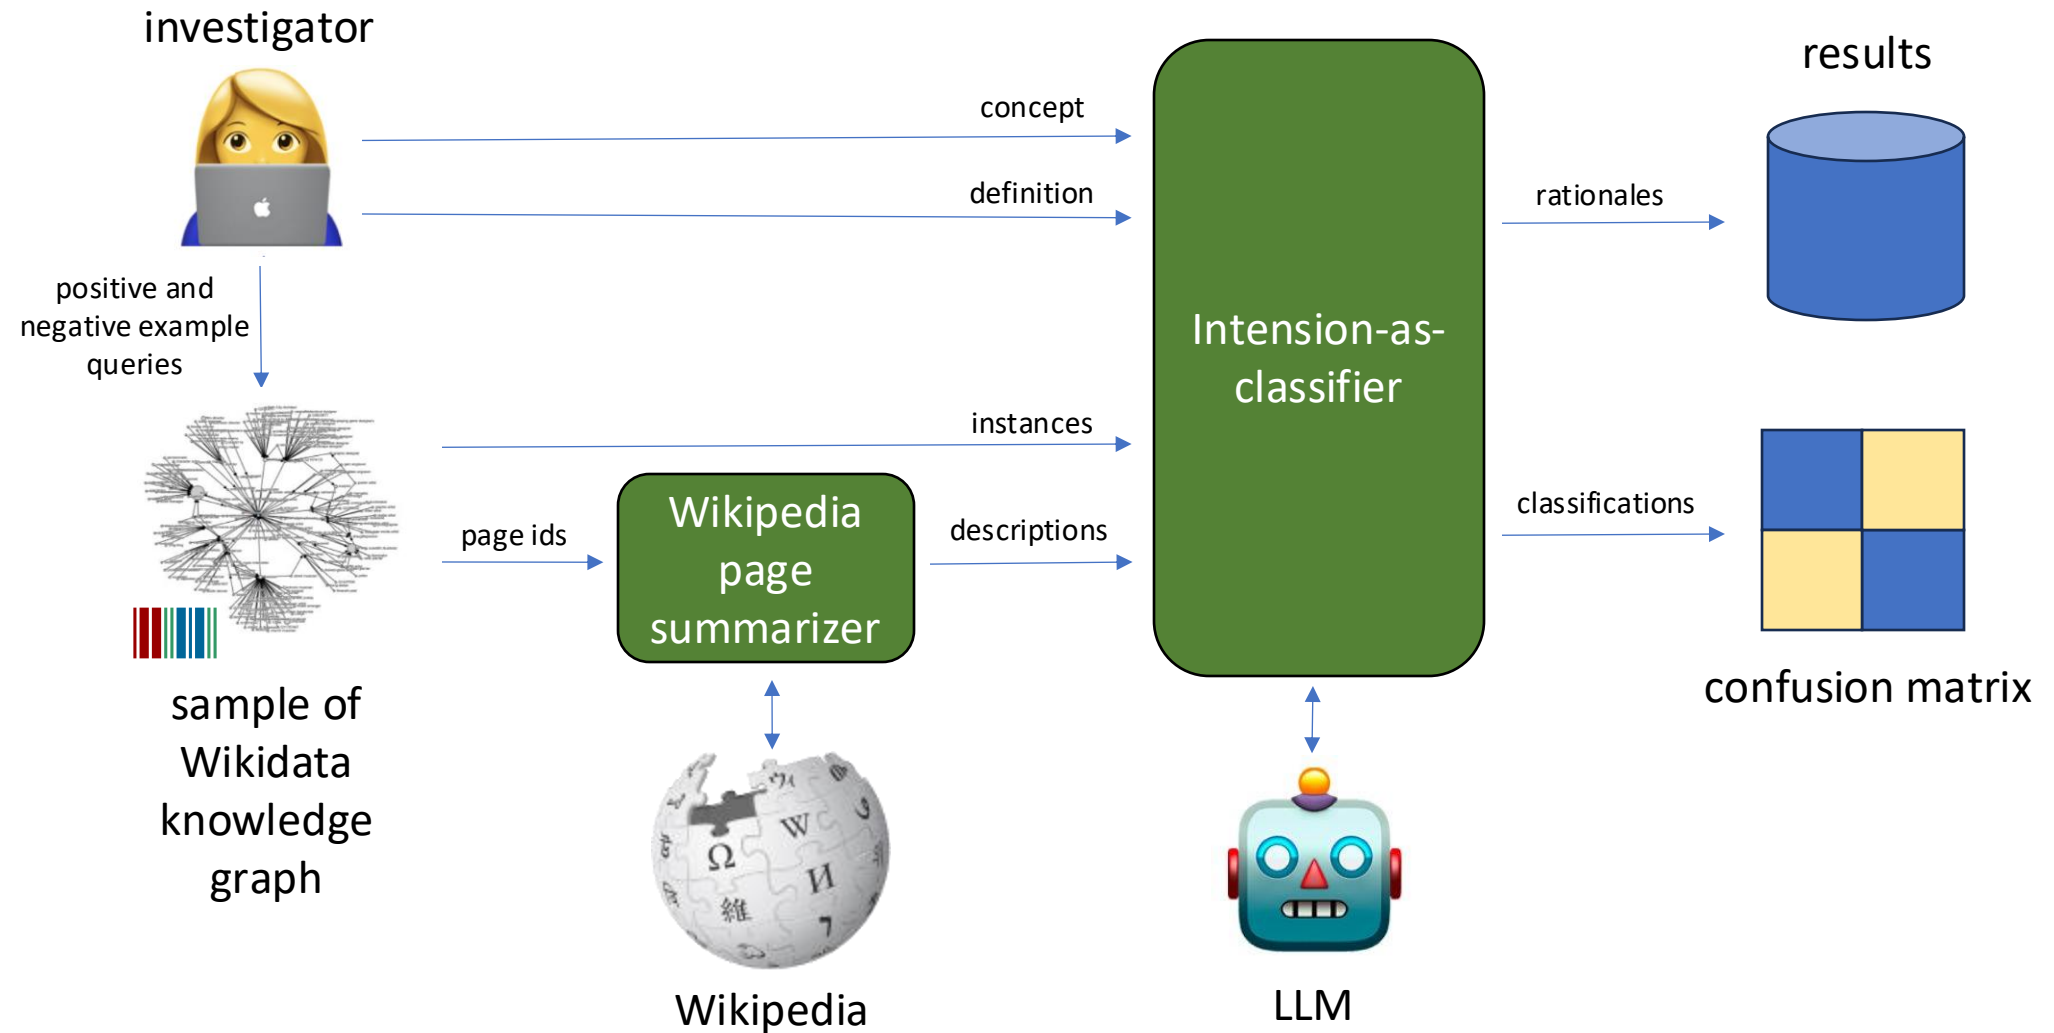
Knowledge graphs

- A *knowledge graph* G is a set of *triples* $\{(s, p, o) \mid s \in S, p \in P, o \in O\}$, where
 - I is a countable set of IRIs (Internationalized Resource Identifiers)
 - B is a countable set of blank nodes
 - L is a countable set of literals
 - $S \subset I \cup B$ is the set of *subjects* in G
 - $P \subset I$ is the set of *properties* in G
 - $O \subset I \cup B \cup L$ is the set of *objects* in G
- Let `instanceOf`, `subClassOf`, `label` $\in P$ denote an instance-of relation, a subclass-of relation, and a label property in G , respectively.
- A *concept* $c \in I \cup B$ is an entity such that $\exists(s, \text{subClassOf}, o) \in G \mid s = c \vee o = c$.

Extensions in knowledge graphs

- The *extension in G* of a concept $c \in G$ is defined recursively, such that:
 - $\mathbf{ext}_G(c) = \bigcup_{i \in \mathbb{N}} \mathbf{ext}_i(c)$
 - $\mathbf{ext}_0(c) = \{e \mid \exists(e, \mathbf{instanceOf}, c) \in G\}$
 - $\mathbf{ext}_{i+1}(c) = \mathbf{ext}_i(c) \cup \{e \mid e \in \mathbf{ext}(c') \wedge \exists(c', \mathbf{subClassOf}, c) \in G\}$



Carnap's Robot Redux (CRR)



Investigations

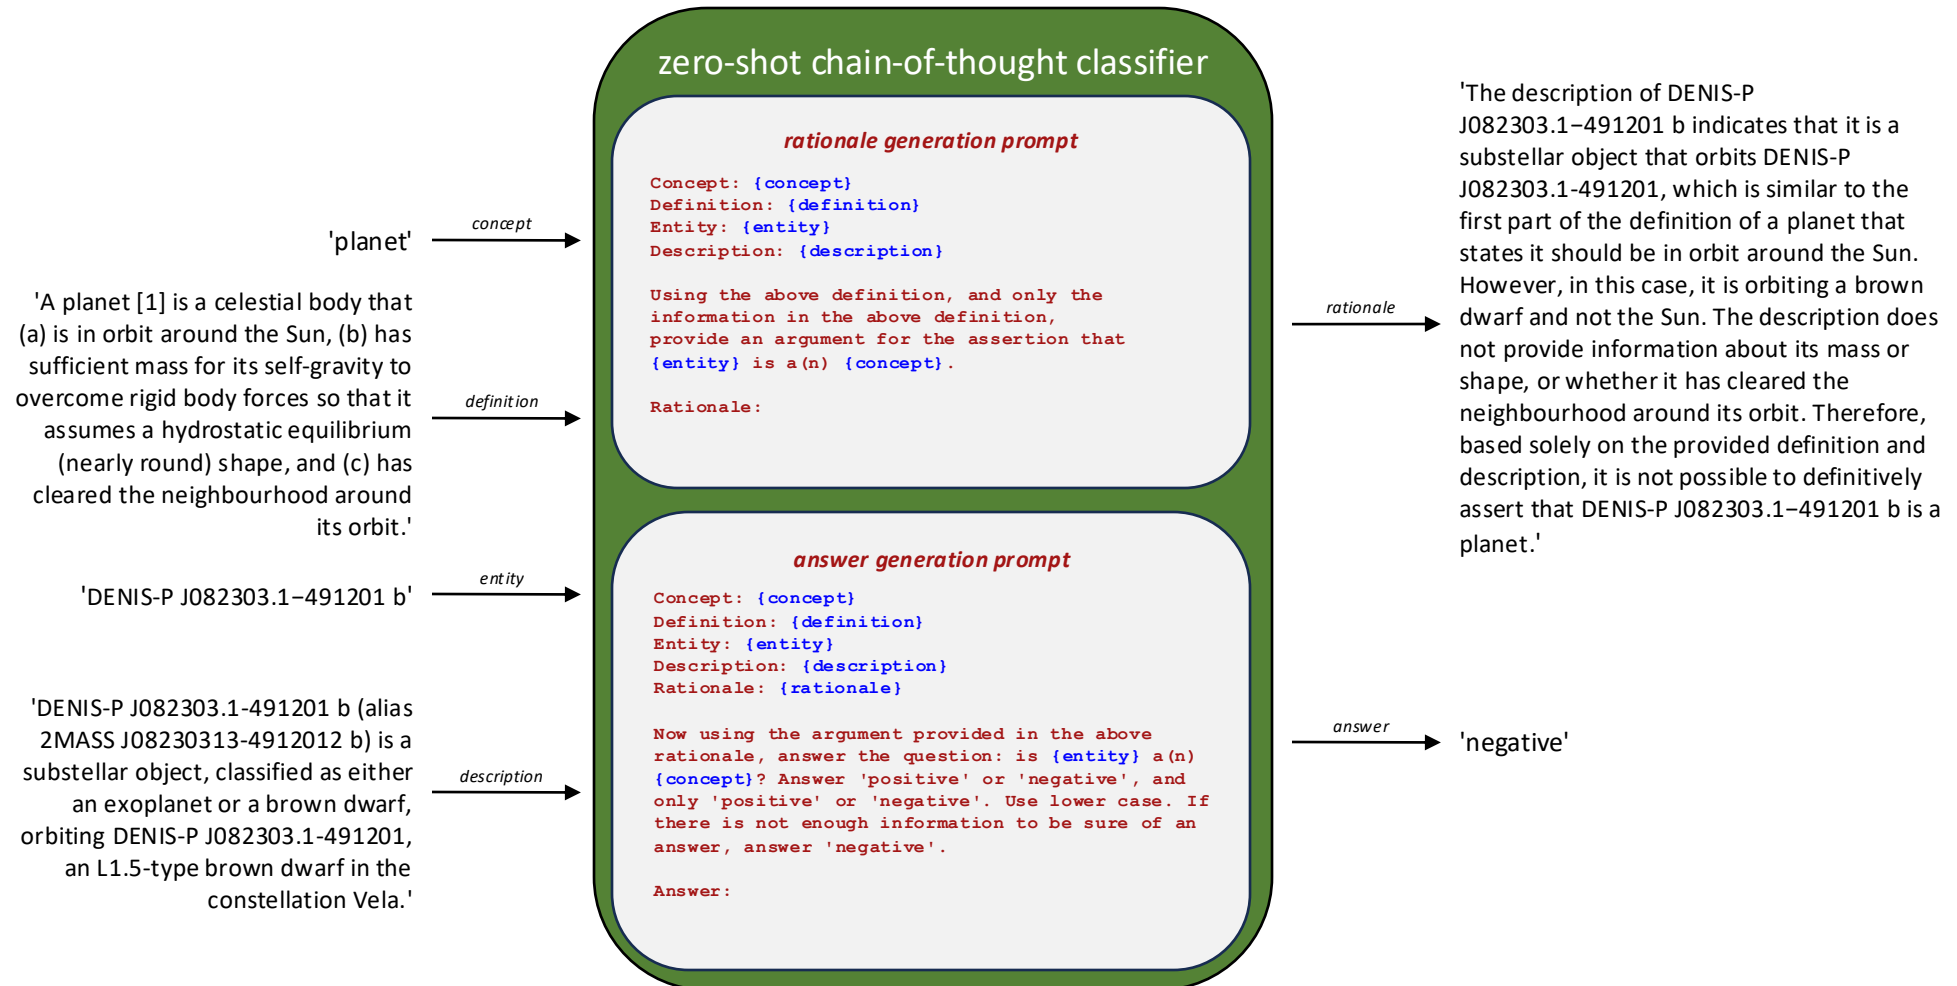
- Our main approach is to analyze intensions and extensions in knowledge graphs using Carnap's Robot Redux
- We conducted two investigations to explore this approach
 - Conceptual engineering: examine intensions related to conceptual engineering projects occurring in Wikidata
 - Knowledge graph refinement: examine intensions of class membership relations more broadly from Wikidata

Classifiers as targets of conceptual engineering



- Conceptual engineering (CE) is a philosophical methodology concerned with “the design, implementation, and evaluation of concepts” (Chalmers 2020)
- The goals of conceptual engineering are varied, e.g.,
 - Achieving greater clarity and precision in argumentation and scientific discourse (Dutilh Novaes & Reck 2017; Justus 2012)
 - Altering terminology to advance the cause of social justice (Haslanger 2000; Manne, 2017; Podosky, 2022)
- An important question for a theory of CE is the nature of its targets, i.e., “what conceptual engineers are (or should be) trying to engineer” (Koch, Löhr & Pinder 2023)
- Nado (2021) proposes as targets *classification procedures*, defined as abstract ‘recipes’ which sort entities “into an ‘in’-group and an ‘out’-group”

Implementing Nado's classification procedures



Experiments

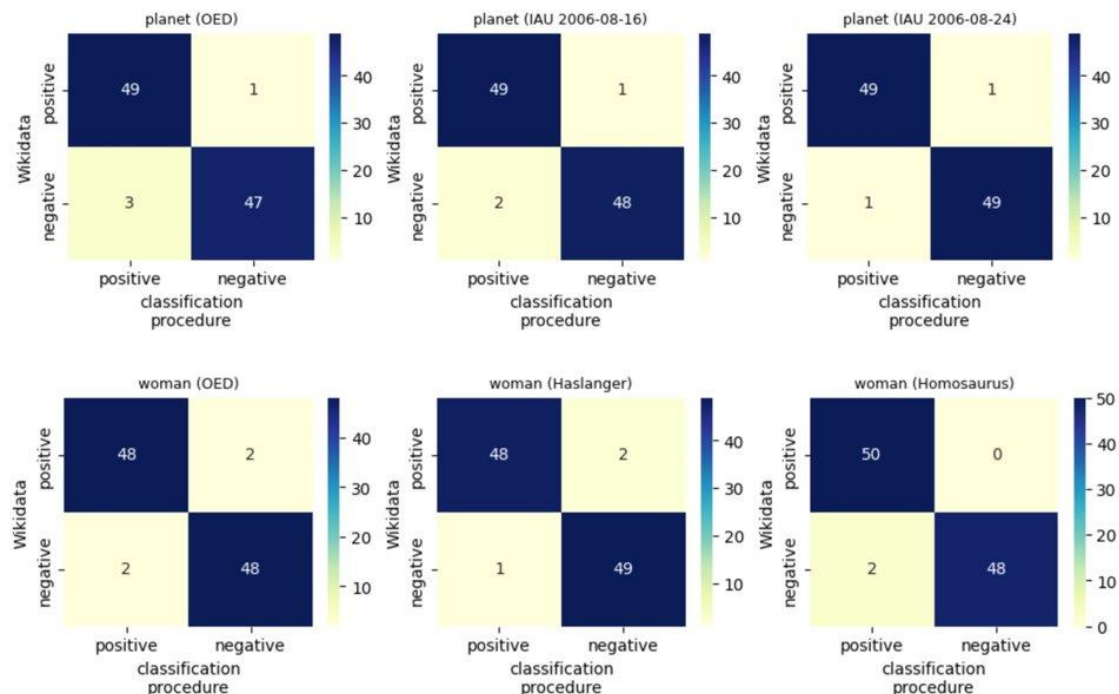
- PLANET

- Definitions: OED, IAU 2006-08-16, IAU 2006-08-24
- Positive examples: sample of 50 instances (P31) of planet (Q634)
- Negative examples: sample of 50 instances of substellar object (Q3132741) that are not instances of planet

- WOMAN

- Definitions: OED, Haslanger, Homosaurus
- Positive examples: sample of 50 instances with sex or gender (P21) either female (Q6581072) or trans woman (Q1052281)
- Negative examples: sample of 50 instances with sex or gender either male (Q6581097), non-binary (Q48270), or trans man (Q2449503)

concept	definition	Cohen's kappa	F1 macro	FN	FP
PLANET	IAU 2006-08-24	0.96	0.98	1	1
	IAU 2006-08-16	0.94	0.97	1	2
	OED	0.92	0.96	1	3
WOMAN	Homosaurus	0.96	0.98	0	2
	Haslanger	0.94	0.97	2	1
	OED	0.92	0.96	2	2



Error analysis: KG failure

```
display_errors(homosaurus_women, homosaurus_women_results)
```

Python

woman (Homosaurus)

Definition

Adults who self-identify as women and understand their gender in terms of Western conceptions of womanness, femaleness, and/or femininity. The term has typically been defined as adult female humans, though not all women identify with the term 'female' depending on the context in which it is used.

False positives

name	description	rationale
Michelle Rojas	Michelle Rojas (née Pennington; born December 16, 1987) is an American voice actress from Dallas, Texas. Affiliated with Funimation and Bang Zoom! Entertainment, some of her major roles are Tohka Yatogami in Date A Live, Zuikaku in Azur Lane, Kanade Sakurada in Castle Town Dandelion, Mikoto Sakuragawa in Gonna be the Twin-Tail!!, Kō Yagami in New Game!, Nashiro Yasuhisa in Tokyo Ghoul, Maya Sato in Classroom of the Elite, Toka Yada in Assassination Classroom, Shion in That Time I Got Reincarnated as a Slime, Roxy Migurdia in Mushoku Tensei: Jobless Reincarnation, and Yamato in One Piece.	The description of Michelle Rojas does not explicitly state her gender identity. However, it does use female pronouns ("her") to refer to her, which typically aligns with a self-identification as a woman. Furthermore, she has taken on roles that are typically associated with women, such as female characters in various shows. Therefore, based on the information provided and the definition of a woman, it can be argued that Michelle Rojas is a woman.
Linden A. Lewis	Linden A. Lewis is a science fiction author, best known for their debut novel The First Sister. Lewis is queer and uses she/they pronouns.	The definition of a woman includes adults who self-identify as women. Linden A. Lewis uses she/they pronouns, which indicates that they may identify as a woman. Therefore, based on the provided definition, it can be argued that Linden A. Lewis is a woman.

Error analysis: LLM failure

woman (Haslanger)

Definition

S is a woman iff (i) S is regularly and for the most part observed or imagined to have certain bodily features presumed to be evidence of a female's biological role in reproduction; (ii) that S has these features marks S within the dominant ideology of S's society as someone who ought to occupy certain kinds of social position that are in fact subordinate (and so motivates and justifies S's occupying such a position); and (iii) the fact that S satisfies (i) and (ii) plays a role in S's systematic subordination, that is, along some dimension, S's social position is oppressive, and S's satisfying (i) and (ii) plays a role in that dimension of subordination

False negatives

name	description	rationale
Waltraud Klasnic	Waltraud Klasnic (née Tschiltsch, born 27 October 1945) is an Austrian politician who was Landeshauptfrau (governor) of Styria from 1996 until 2005.	The definition of a woman provided includes three main criteria: (i) having certain bodily features associated with the female biological role in reproduction, (ii) being marked within the dominant ideology of one's society as someone who should occupy certain subordinate social positions, and (iii) the fact that satisfying (i) and (ii) contributes to one's systematic subordination. In the case of Waltraud Klasnic, while the description does not explicitly state her biological features or her societal role, it does mention that she is an Austrian politician who served as the governor of Styria. This implies that she is recognized and accepted in a significant societal role, which suggests that she likely satisfies criteria (ii). However, the description does not provide enough information to definitively assert that Waltraud Klasnic satisfies all the criteria in the definition of a woman. For instance, it does not provide information about her biological features or whether she experiences systematic subordination due to her gender. Therefore, based solely on the provided definition and description, we cannot conclusively argue that Waltraud Klasnic is a woman.

Knowledge graph refinement

- Knowledge graph (KG) refinement is an increasingly important topic
 - KGs can be messy (e.g., Wikidata quality issues)
 - Human evaluation is expensive
- CRR can be used to support the evaluation of *class membership relations* in a KG
 - Class membership relations are a principal way in which KGs represent classification schemes
 - Classification schemes are a crucial aspect of knowledge infrastructure, often with significant implications for social policy and scientific consensus
 - Class membership is also a key concern in data governance, which depends on "providing a set of mappings from a representation language to agreed-upon concepts in the real world" (Khatri & Brown 2010)

Experiments

- KGs: Wikidata, CaLiGraph
- Data:
 - 20 Wikidata classes, 19 CaLiGraph classes randomly sampled
 - 20 positive and ~20 negative examples randomly sampled per class
 - Serializations limited to 20 triples per example
 - ~800 class/entity pairs per KG
- Error analysis: manually review FNs, FPs with rationales and assign error to LLM or KG
 - LLM errors: incorrect reasoning, missing data
 - KG errors: missing relation, incorrect relation
 - Error analysis performed for gpt-4-0125-preview
- Results: classifiers can exhibit good alignment with KGs
 - One LLM was in moderate agreement ($\kappa > 0.60$) with Wikidata
 - Four were in moderate agreement with CaLiGraph

KG	LLM	ACC	AUC	F1	κ
Wikidata	gpt-4-0125-preview	0.830	0.830	0.823	0.660
	gemma-7b-it	0.726	0.727	0.705	0.454
	Mixtral-8x7B-Instruct-v0.1	0.697	0.696	0.654	0.393
	Mistral-7B-Instruct-v0.2	0.671	0.671	0.620	0.342
	gemma-2b-it	0.674	0.670	0.629	0.330
	gpt-3.5-turbo	0.627	0.627	0.547	0.255
	Llama-2-70b-chat-hf	0.631	0.616	0.569	0.239
CaLiGraph	gpt-4-0125-preview	0.900	0.893	0.889	0.788
	Mixtral-8x7B-Instruct-v0.1	0.893	0.884	0.874	0.767
	gpt-3.5-turbo	0.842	0.833	0.815	0.665
	Mistral-7B-Instruct-v0.2	0.812	0.803	0.779	0.605
	gemma-7b-it	0.783	0.774	0.750	0.547
	Llama-2-70b-chat-hf	0.637	0.625	0.558	0.252
	gemma-2b-it	0.563	0.543	0.422	0.090

Error analysis of gpt-4-0125-preview results

KG	N	FP	FN	human-KG κ	human-LLM κ	missing data	missing relation	incorrect relation	incorrect reasoning
Wikidata	136	46	90	0.243	-0.241	34 (25.0%)	15 (11.0%)	33 (24.3%)	54 (39.7%)
CaLiGraph	77	27	50	-0.295	0.198	28 (36.4%)	19 (24.7%)	20 (26.0%)	10 (13.0%)
	213	73	140			62 (29.1%)	34 (16.0%)	53 (24.9%)	64 (30.0%)

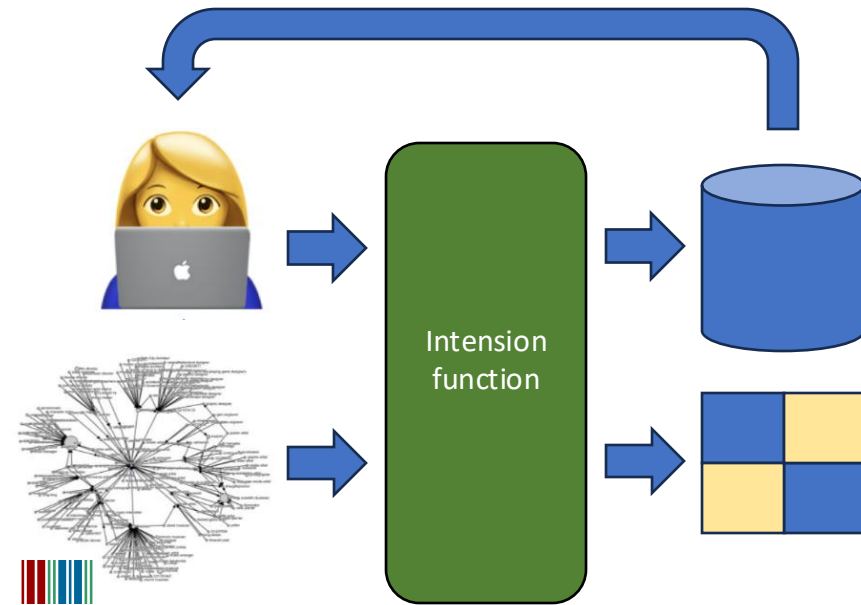
- Classifiers can detect missing or incorrect relations
 - 40.9% of errors were due to the problems with the KG
 - 29.1% of errors were due to missing or insufficient data in the entity description
 - 30.0% of errors due to incorrect reasoning by the LLM
- Pairwise human-KG and human-LLM agreement differed between the KGs
 - Human showed fair agreement with Wikidata and no agreement with the classifier
 - Human showed slight agreement with the classifier and no agreement with CaLiGraph

Objections

- Trustworthiness
 - Humans can be misled by rationales produced by LLM classifier
 - "Philosophers are (usually) competent natural language speakers and especially keen to subtle differences in meaning." (Justus 2012, p. 172)
- Groundedness
 - LLMs may or may not be able to make meaningful statements
 - Nado's Practical Role Account (Nado 2023): "success in conceptual engineering is a matter of devising a tool that will effectively fulfil an intended practical role"


CRR as a tool for experimental conceptual engineering

- Carnap (1955) "obviously resembles the survey methodology of x-phi—indeed it warrants identifying Carnap as an early pioneer of x-phi" (Shepherd & Justus 2015)
 - Explication preparation as a facet of experimental conceptual engineering
- CRR might arguably be viewed as an instance of a corpus method for evidential x-phi (Fischer & Sytsma 2022)
- Opinions differ as to whether LLMs can (Dillon et al. 2023) or cannot (Harding et al. 2024) be used as a substitute for surveying humans' opinion or simulating folk belief, but in general it is accepted that they can be useful for preparatory exploration
 - This again is a take consistent with Nado's Practical Role Account

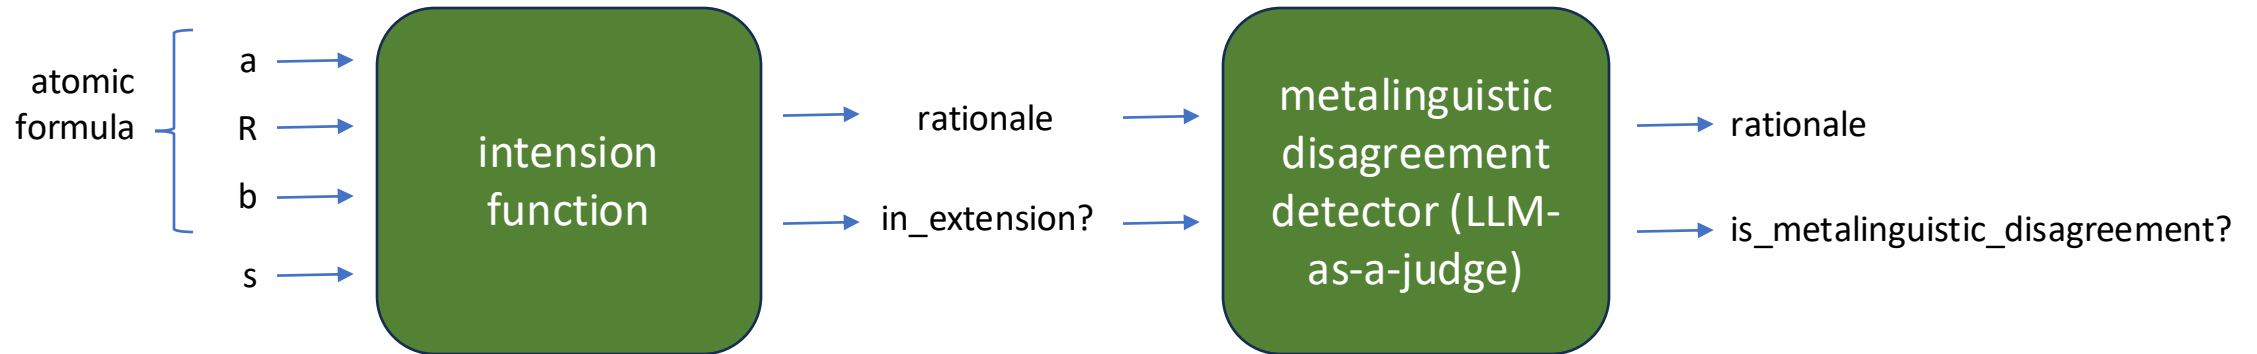


CRR as a tool for addressing the implementation problem in conceptual engineering

- The implementation problem in conceptual engineering poses the question of whether (re)engineered concepts can be effectively adopted by a population of human speakers (Cappelen, 2018; Jorem, 2021)
- Knowledge graphs have a direct and material impact on society by virtue of their use in online search, discovery, and recommendation.
- Using CCR, we can guide ameliorative changes to facts in a knowledge base to better align with a proposed definition, shifting the semantic meaning of a concept for the specific linguistic subgroup (Matsui, 2024) constituted by users of such online knowledge bases.



Recent work: Do LLMs exhibit metalinguistic disagreement with knowledge graphs?



- Sample of 250 alignments between Wikidata triples and Wikipedia abstracts from T-REx dataset
- An intension-as-a-classifier was executed to obtain a rationale and a truth value for each of the sampled triples and aligned abstracts
- Each classification was then processed by a second zero-shot chain-of-thought classifier (using gpt-4o-2024-05-13) that classified whether the truth-value-assigning classifier's rationale indicated a metalinguistic disagreement

LLM	FN	FNR	metalinguistic disagreements	MDR
gpt-4o-2024-05-13	26	0.104	10	0.040
gpt-4-0125-preview	33	0.132	16	0.064
claude-3-haiku-20240307	42	0.168	11	0.044
claude-3-opus-20240229	51	0.204	14	0.056
claude-3-5-sonnet-20240620	52	0.208	16	0.064
meta-llama/Meta-Llama-3-70b-Instruct	71	0.284	28	0.112
gpt-3.5-turbo	75	0.300	29	0.116
mistralai/Mixtral-8x7B-Instruct-v0.1	78	0.312	29	0.116
mistralai/Mistral-7B-Instruct-v0.3	126	0.504	66	0.264

Example rationales exhibiting metalinguistic disagreement

Wikidata triple	rationale for negative finding
James II of England (Q126188) followed by (P156) Mary II of England (Q130812)	Based on the information provided in the hypothetical world, the knowledge graph triple “Subject: James II of England, Predicate: followed by, Object: Mary II of England” is false. The passage states that after James II was deposed in 1688, he was replaced by his daughter Mary II, ruling jointly with her husband William III. This means that Mary II did not immediately follow James II, but rather replaced him as the monarch. The predicate “followed by” implies that Mary II succeeded James II directly, which is not the case.
chocolate (Q195) made from material (P186) sugar (Q11002)	Based on the information provided, the knowledge graph triple “Subject: chocolate, Predicate: made from material, Object: sugar” is false. The passage clearly states that chocolate is made from the seeds of the cacao tree, specifically Theobroma cacao, and that it is roasted and ground to produce the chocolate. There is no mention of sugar being a material that chocolate is made from. In fact, the passage states that sugar is added to chocolate to create sweet chocolate, but it is not a primary ingredient that chocolate is made from. Therefore, the given knowledge graph triple is false.

Future work

- How pervasive is metalinguistic disagreement between LLMs and knowledge bases?
- How trustworthy are rationales as propositional attitude reports?
- What is the relationship between hyperintensionality, topicality, and metalinguistic disagreement?
- Can we extend our LLM-based intensional semantics to a two-component semantics?
- Can this yield a logic of belief for LLM agents?

Conclusions

- We've shown how LLMs can be used to turn Carnap's Robot into a cognitive tool for working with intensions
 - For purposes of both conceptual engineering and knowledge engineering
- Carnap's Robot, a thought experiment from 70 years ago, turns out to be a useful tool for exploring the factuality and trustworthiness of LLMs
 - Carnap anticipated the central challenge associated with working with LLMs: can we use them as black boxes, or must we solve the problem of mechanistic interpretability?
- We would love to answer the questions around belief and meaning in LLMs, but we do not have to definitively answer those questions for them to be useful
 - LLMs open a path towards experimental conceptual engineering
 - LLMs used for knowledge base refinement can provide a way to address the implementation problem in conceptual engineering



UNIVERSITEIT
VAN AMSTERDAM

INDE lab

Thank you!

b.p.allen@uva.nl

<https://www.bradleypallen.org>